



Referrals Decision Guide

Companion to the [Referrals](#) docs.

A referral needs more than a send button. Use this guide with your clinical and operations teams to decide who owns each step, what the receiving team needs, and how you will know the loop is closed.

Section 1: Use Case & Participants

1.1 Which side(s) of the referral lifecycle do you build for?

- Outbound only – you originate referrals to others
- Inbound only – you receive referrals from others
- Both

Why: determines which Section 3 deep dives apply.

1.2 Who interacts with referrals, and is the patient part of the workflow?

- Sending side: which roles author, edit, and send referrals (clinicians, MAs, coordinators, intake staff)?
- Receiving side: which roles triage, accept, and work referrals on intake?
- Patient: do patients see referrals, get notified, or take action (e.g. view in portal, self-schedule, decline)?

Why: drives how forms are designed (3.2), who notes are written for (3.4), who owns the work (3.7), and whether patient-facing content is in scope.

1.3 Where do receiving providers live relative to your tenant?

- Inside your tenant only
- Partner tenants or organizations
- External organizations without FHIR connectivity
- Some combination

Why: drives how recipients are stored and selected (3.3), and which outbound channels you support (3.5).

For teams sharing one project, create receiving work against the shared ServiceRequest directly. Add transmission and identifier reconciliation when records cross systems. See [internal and external referrals](#).

1.4 Is this a new build or replacing/extending something, and what other systems must it work with?

- If replacing: what system or workflow, and what gaps drove the change?
- If new: is referrals central to the product or additive?
- What other systems will referrals flow to or from (source EHR, downstream EHRs at receivers, billing, HIE, payer portals, include DirectMessage)?

Why: surfaces migration needs and integration boundaries that affect outbound channels (3.5), inbound intake (3.6), and matching of patient and provider records across systems.

Section 2: Feature Scoping

Go through each row together. For each cell under Yes / No / Nice-to-have / Not sure, mark the customer's answer. The \$ column points to the deep-dive subsection that covers each feature.

#	Feature	\$	Yes	No	Nice-to-have	Not sure
1	Multiple referral forms by specialty (vs one generic form)	3.2				
2	Tie every referral to the visit where it was created	3.2				
3	Pull in chart context automatically when sending (recent labs, problems, meds, allergies, coverage)	3.4				
4	Different audiences for notes on the referral – patient, receiving provider, internal staff, or mix	3.4				
5	Maintain your own directory of referral recipients (clinicians, sites, queues) in Medplum, vs sourcing from another system	3.3				

#	Feature	\$	Yes	No	Nice-to-have	Not sure
6	Send referrals to recipients outside your product (printed or PDF summary, care-summary document, fax, secure messaging, email)	3.5				
7	Accept referrals from outside (fax or scanned mail, secure inbox, Direct Message, care-summary drops, API or health-information exchange, typed in by staff)	3.6				
8	Route / assign referrals and track work on open referrals (intake pools, ownership, aging or SLA dashboards)	3.7				
9	Prior authorization workflow (eligibility checks, payer auth submissions before scheduling)	3.7				
10	In-network / coverage check against patient insurance when selecting or verifying a recipient	3.3 3.7				
11	Custom status steps for where the referral sits in workflow (e.g. sent → accepted → scheduled → done)	3.8				
12	Tie results and visit notes back to the original referral so the thread stays complete	3.9				

Section 3: Feature Deep Dives

Cover each feature flagged Yes or Nice-to-have in Section 2. The goal is to land on a clear recommended approach by the end of each section.

Section 3 is grouped into four lanes:

Lane	Subsections	When to read
Referral Resource	3.1	Always

Lane	Subsections	When to read
Sending Referrals	3.2 – 3.5	When you originate referrals
Receiving Referrals	3.6	When you accept referrals from outside
Lifecycle & Processing	3.7 – 3.9	Always

Referral Resource

3.1 Referral Semantics – Type, Coding, Intent, Priority

Each independently managed requested service uses a ServiceRequest; requests authorized together can share a requisition identifier. Decide how you classify it, how urgency is expressed, whether you expose a staff-facing reference number, and whether drafts exist before send.

Questions:

- How do you classify referrals (specialty, service, program), using standard codes, your own list, or both?
- What urgency options do users pick, and should they match triage and reporting?
- Do you show a referral number or reference beyond the system id?
- Is the referral final at send/commit, or do you need a draft first?

Situation	Approach
Category and service type	Bind category / code to agreed ValueSets (SNOMED, custom, or mix); drives search and reporting.
Priority	Set ServiceRequest.priority; keep labels aligned with clinical operations.
Display referral number	ServiceRequest.identifier with your namespace; assign at create when needed for ops or correspondence.
Released order	Set intent = order and status = active through the authorized release path. Sending is tracked separately from order status.
Draft before release	Use status = draft with intent = order for an order being prepared. For a true proposal, retain it and create a separate authorized order with basedOn referencing the proposal.

Sending Referrals

3.2 Capture Flow – Form Library & Originating Context

Decide how referral intake is structured in the UI – one flow vs several by specialty or program, how answers are constrained, and whether every referral must link to an originating visit.

Questions:

- Do you use one referral intake flow for everyone, or different flows by specialty, service, or program?
- Which fields must be chosen from lists or structured inputs (e.g. specialty, recipient, urgency), which allow free text, and what is required before the referral can be sent?
- Must every referral tie to the visit where it was created, or can referrals start without a linked visit (e.g. intake, triage, outreach)?

Situation	Approach
One form	Use one Questionnaire and retain its QuestionnaireResponse. Choose one extraction strategy for the form: SDC template extraction or a Bot mapping to ServiceRequest.
Specialty forms (e.g., variation based on specialty, service, etc.) or supplemental forms (specific extra information needed based on receiving clinic preferences)	Multiple Questionnaires; shared core fields + specialty-specific items (specialty forms) or additional Questionnaires specific to receiving clinics (supplemental forms); map each to ServiceRequest.
Encounter required	Require encounter when referrals always start from a scheduled or in-progress visit.
Encounter optional	Allow missing encounter when referrals can start outside a visit (e.g. intake); enforce subject and clinical reason instead.

3.3 Recipient & Directory Model

Decide who users may send a referral *to* (person vs site vs queue) and where your recipient directory lives – maintained in Medplum, fed from another system, or both.

⚠ **One-way door:** If recipients start as free text and you move to a structured directory later, expect to reconcile historical referrals against the new directory.

Questions:

- Can a referral go to a specific clinician, to a site, department, or queue, or either, depending on referral type?

- Is your recipient directory edited and stored in Medplum, synced from another system, or mixed?
- Should the recipient picker show or filter by in-network status for the patient's coverage – and where does that network status come from (synced payer provider directory, internal preferred-provider list, or payer-specific eligibility evidence)?

Situation	Approach
Named individual	Use performer → PractitionerRole when the practice/service context matters, or Practitioner for the individual alone.
Pool or department	performer → Organization or HealthcareService when the referral is to a group or queue.
Directory in Medplum	Maintain Practitioner, PractitionerRole, Organization, and HealthcareService; use Location for physical sites and reference-based recipient pickers.
External directory	Sync or resolve selections to FHIR references at submit time via Bot or integration.
Network-aware picker	Use authoritative payer directory or contract data for the provider, service, location, and effective date. Record unknown/stale status explicitly. A preferred-provider list or eligibility response alone does not establish network participation.

3.4 What the Referral Carries – Notes & Clinical Package

Decide what leaves with the referral, how much is automated versus curated at send, and which audiences see each part of the narrative.

Questions:

- What clinical material must travel with the referral (reason for referral, problems or diagnoses, key meds and allergies, recent labs or imaging, coverage), and what should be filled in automatically versus selected or edited by the sender before send?
- Should internal staff-only notes appear on the referral, or only content intended for the receiving provider and the patient?

Situation	Approach
Receiver + patient messaging	Use note for shared clinical narrative and patientInstruction for patient-facing instructions. These fields do not enforce different audiences; keep internal notes in separately protected resources and exclude them from the package.

Situation	Approach
Lean vs rich package	Link the clinical reason through reasonReference or reasonCode; use supportingInfo for selected evidence. Store files in Binary and link them through DocumentReference attachment URLs, without inline base64.
Context at send	Let the sender review selected chart context. Retain the actual transmitted document or payload and its source versions so later chart edits do not change the send history.

3.5 Outbound Transmission – Channels & Payload Format

Decide how referrals leave your product – whether recipients get structured exchange versus documents – and what you produce for each sending path.

Questions:

- Do any recipients accept referrals through structured exchange with your system, or only as documents (e.g. PDF or care-summary packages)?
- What outbound channels do you plan to use? (e.g., secure email, fax, API?)
- For each outbound channel you plan to use, what do you actually send or generate – structured data export, PDF, care-summary document, or a mix?
- Is the outbound channel the same for every recipient (e.g. always fax) or does it vary per recipient – and if it varies, where does that preference live (recipient directory, referral type, picked at send time)?
- For each channel, do you need to confirm delivery, retry on failure, and show send history to staff – and what counts as “successfully sent”?

Situation	Approach
FHIR-capable peer	Send or expose ServiceRequest and related resources via API; record transmission with Communication when useful for audit.
No FHIR API (e.g., secure email, fax channels)	Generate a reviewed PDF, or use a validated C-CDA generator matching the exchange contract. Store the result as DocumentReference/Binary and link the sent package from Communication.
More than one channel	Keep one ServiceRequest as source of truth; run separate send paths per channel without divergent clinical content.

Receiving Referrals

3.6 Inbound Intake – Channels & Ingestion

Decide how inbound referrals get created, reviewed, and matched to patients.

Questions:

- Which channels feed referrals into your product (e.g. automated exchange with a partner, fax or scanned documents, imported files or care-summary packages, typed in by staff)?
- For each channel, does a referral become active in your system automatically, or only after staff review or cleanup?
- When inbound patient information is incomplete or inconsistent, how do you find or confirm the correct patient record?
- After accepting (or declining) a new inbound referral, do you need to send an acknowledgment or response back to the originator – and through what path?
- When the same referral arrives more than once (e.g. fax + secure inbox), how do you detect and handle duplicates?

Situation	Approach
Structured FHIR inbound	Validate ServiceRequest and related resources, resolve patient identity and local references, and preserve source identifiers. Track receiving review and acceptance with a Task.
Structured documents or payloads you can parse (e.g. C-CDA XML, FHIR Bundles)	Bot or integration parses and maps fields → ServiceRequest (+ related resources); retain source as DocumentReference / Binary. Native structured interchange does not require OCR; image-only or bitmap PDF renditions may.
Fax, scanned, or image-like documents	Retain the Binary and DocumentReference; extract fields with OCR and review uncertain identity or clinical content. Use a restricted intake Task until the patient is confirmed, then create or reconcile the ServiceRequest.
Patient identity	Match trusted Patient identifiers with their namespaces and route ambiguity for review. Deduplicate referrals separately using the source referral identifier and conditional create.

Lifecycle & Processing

3.7 Working the Referral – Assignment, Authorization, Scheduling

After a referral is created or accepted, decide how it gets to the right person, what authorization and coverage steps gate the work, and how it turns into an actual visit or service that ties back to the referral.

Note: Scheduling is covered in the [Scheduling Decision Guide](#). This section stays at the level needed to align scheduling and authorization steps with the referral lifecycle.

Questions:

- Who picks up a new referral first – an intake pool, a department queue, or an assigned coordinator (+ how does assignment logic work) – and how is ownership tracked and reassigned over time?
- Does the referral need to confirm in-network status or covered benefits for the patient's plan before scheduling – and is that check separate from prior auth, or rolled into the same eligibility request?
- Does the referral need to clear an authorization or coverage check before scheduling, and how do you track that authorization through to approval or denial?
- When the referral becomes an actual visit (appointment, procedure, completed encounter), how do you tie that visit back to the originating referral?
- Do you track aging or service-level expectations on open referrals (e.g. days since received, days pending an action), and how do staff find work that is overdue?
- What happens after unsuccessful outreach, a patient decline, a no-show, or expiration? Who can stop attempts, and how is the referrer informed?
- Which referral and authorization details must be available to billing, including effective dates and visit or unit limits?

Situation	Approach
Ownership and assignment	Set Task.for to the patient and focus to the ServiceRequest. Use owner for the accountable person/group and performerType for required roles; claim and reassign with version checks.
Check in-network status / covered benefits	Use CoverageEligibilityRequest with purpose: benefits where supported. Evaluate network participation against authoritative payer data for the selected provider/service; do not assume a universal eligibility or clearinghouse response.
Check whether prior auth is required	Use CoverageEligibilityRequest with purpose: auth-requirements where supported. Interpret authorizationRequired in the response's insurance items and supporting requirements under the payer contract.

Situation	Approach
Submit a prior auth request	Use Claim with use: preauthorization and interpret ClaimResponse under the chosen payer API, implementation guide, or EDI mapping. Configure that integration explicitly; transport success does not mean approval.
Authorization gate	Keep the Task on hold with an authorization-pending business stage until the required approval evidence is present and valid for the service/date. ClaimResponse.outcome: complete alone does not prove approval.
Scheduling the work	See the Scheduling Decision Guide . From the referral side, link the resulting Appointment and downstream Encounter back to the referral via basedOn → ServiceRequest.
Unsuccessful scheduling	Record outreach against the referral, assign retry/stop decisions, and close each work item according to its scope. Report the disposition to the referring team; see unsuccessful scheduling .
Billing handoff	When required, use Claim.referral for the ServiceRequest and Claim.insurance.preAuthRef for the authorization number. Preserve payer-specific validity and service limits.
Aging and SLA	Use Task history or explicit transitions for time in stage; lastModified may reflect unrelated edits. Follow the Task planned-date convention , range-aware search, and a defined pause policy.

3.8 Lifecycle & BusinessStatus CodeSystem

Decide how you track a referral from request through completion – or cancellation – including naming those steps and handling cancel vs reroute.

Questions:

- Which milestones matter enough to show and operate on (e.g. sent, accepted, scheduled, declined, unable to schedule, completed, cancelled)? Which updates must reach the referrer?
- How do you standardize labels for those milestones so reports, queues, and integrations all use the same meanings?

- If a referral is cancelled or the recipient changes, do you revise the existing referral or create a new one?

Situation	Approach
Linear milestones	Use R4 Task.status for lifecycle and governed local businessStatus codes for workflow stages. Define the scope and completion evidence of each Task.
External signals	Correlate and deduplicate callbacks before updating work. Delivery, receiving-team acceptance, service completion, and returned-result review are separate events.
Cancel / reroute	For an unchanged clinical order, the authorized workflow can update performer and resolve old Tasks. For a replacement order, use a new ServiceRequest with replaces and end the old request appropriately. Preserve transmission history.

3.9 Results & Closure Round-Trip

Decide what should come back after the referral is acted on, how those items attach to the right referral, and what counts as done for the referring side, the receiving side, or both.

Questions:

- After the referral goes out, what do you expect back – structured results (e.g. labs or imaging), a consult or visit summary, documents or letters, several of those, or nothing beyond tracking the referral itself?
- When results or documents arrive, how do you attach them to the correct referral – including when human confirmation is needed?
- What marks the referral complete for your operations – the referring organization, the receiving one, both, or different definitions for each?

Situation	Approach
Labs / imaging	DiagnosticReport with basedOn → ServiceRequest.
Completed specialist visit	Encounter with basedOn → ServiceRequest when that models your workflow.
Narrative / PDF reply	Use DocumentReference.context.related → ServiceRequest and an attachment URL → Binary. context.encounter is for Encounter or EpisodeOfCare.
Linking	Manual selection or automated match using identifiers you control – define when human confirmation is required.

Situation	Approach
Closure	End each Task with the appropriate status and disposition. Reconcile ServiceRequest completion from service evidence through the responsible authority. For closed-loop care, separately require returned-result review and assignment of follow-up.

Implementation Guides

Use the decisions above to choose the next implementation guide:

- [Creating and Capturing Referrals](#)
- [Recipients and Referral Packages](#)
- [Sending Referrals and Tracking Delivery](#)
- [Receiving and Triaging Referrals](#)
- [Processing and Coordinating Referrals](#)
- [Results and Closing the Loop](#)
- [Worked Referral Example](#)

For assignment, delegation, and longitudinal workflows shared across the practice, use the [Care Coordination Decision Guide](#).